

```

import numpy as np

import torch
import torch.nn as nn
import torchvision
import torchvision.transforms as T
import torch.optim as optim
from torch.utils.data import sampler

import PIL

NOISE_DIM = 96

dtype = torch.cuda.FloatTensor if torch.cuda.is_available() else torch.FloatTensor

def sample_noise(batch_size, dim, seed=None):
    """
    Generate a PyTorch Tensor of uniform random noise.

    Input:
    - batch_size: Integer giving the batch size of noise to generate.
    - dim: Integer giving the dimension of noise to generate.

    Output:
    - A PyTorch Tensor of shape (batch_size, dim) containing uniform
      random noise in the range (-1, 1).
    """
    if seed is not None:
        torch.manual_seed(seed)

    # *****START OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****

    return 2 * torch.rand(batch_size, dim) - 1

    # *****END OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****

def discriminator():
    """
    Build and return a PyTorch model implementing the architecture above.
    """
    # *****START OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****
    model = nn.Sequential(
        Flatten(),
        nn.Linear(28*28, 256),
        nn.LeakyReLU(0.01),
        nn.Linear(256, 256),
        nn.LeakyReLU(0.01),
        nn.Linear(256, 1)
    )
    return model
    # *****END OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****

def generator(noise_dim=NOISE_DIM):
    """
    Build and return a PyTorch model implementing the architecture above.
    """
    # *****START OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****
    model = nn.Sequential(
        nn.Linear(noise_dim, 1024),
        nn.ReLU(True),
        nn.Linear(1024, 1024),
        nn.ReLU(True),
        nn.Linear(1024, 784),
        nn.Tanh()
    )
    return model

```

```
# *****END OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****
```

```
def bce_loss(input, target):  
    """  
    Numerically stable version of the binary cross-entropy loss function in PyTorch.  
  
    Inputs:  
    - input: PyTorch Tensor of shape (N, 1) giving scores.  
    - target: PyTorch Tensor of shape (N, 1) containing 0 and 1 giving targets.  
      dtype is float! (a global dtype is defined above).  
  
    Returns:  
    - A PyTorch Tensor containing the mean BCE loss over the minibatch of input data.  
    """  
    bce = nn.BCEWithLogitsLoss()  
    return bce(input, target)
```

```
def discriminator_loss(logits_real, logits_fake):  
    """  
    Computes the discriminator loss described above.  
  
    Inputs:  
    - logits_real: PyTorch Tensor of shape (N, 1) giving scores for the real data.  
    - logits_fake: PyTorch Tensor of shape (N, 1) giving scores for the fake data.  
  
    Returns:  
    - loss: PyTorch Tensor containing (scalar) the loss for the discriminator.  
    """  
    loss = None  
    # *****START OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****  
  
    labels_real = torch.ones_like(logits_real)  
    labels_fake = torch.zeros_like(logits_fake)  
    loss_real = bce_loss(logits_real, labels_real)  
    loss_fake = bce_loss(logits_fake, labels_fake)  
    return loss_real + loss_fake  
  
    # *****END OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****  
    return loss
```

```
def generator_loss(logits_fake):  
    """  
    Computes the generator loss described above.  
  
    Inputs:  
    - logits_fake: PyTorch Tensor of shape (N, 1) giving scores for the fake data.  
  
    Returns:  
    - loss: PyTorch Tensor containing the (scalar) loss for the generator.  
    """  
    loss = None  
    # *****START OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****  
    return bce_loss(logits_fake, torch.ones_like(logits_fake))  
  
    # *****END OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****  
    return loss
```

```
def get_optimizer(model):  
    """  
    Construct and return an Adam optimizer for the model with learning rate 1e-3,  
    beta1=0.5, and beta2=0.999.  
  
    Input:  
    - model: A PyTorch model that we want to optimize.
```



```

nn.Conv2d(32, 64, kernel_size=5, stride=1),
nn.LeakyReLU(0.01),
nn.MaxPool2d(2, 2),
nn.Flatten(),
nn.Linear(1024, 1024),
nn.LeakyReLU(0.01),
nn.Linear(1024, 1)
)

# *****END OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****
#####
#                                     END OF YOUR CODE                                     #
#####

```

```
def build_dc_generator(noise_dim=NOISE_DIM):
```

```
    """
```

```
    Build and return a PyTorch model implementing the DCGAN generator using
    the architecture described above.
    """
```

```

#####
# TODO: Implement architecture                                     #
#                                                                 #
# HINT: nn.Sequential might be helpful.                         #
#####
# *****START OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****

```

```

return nn.Sequential(
    nn.Linear(noise_dim, 1024),
    nn.ReLU(True),
    nn.BatchNorm1d(1024),

    nn.Linear(1024, 128 * 7 * 7),
    nn.ReLU(True),
    nn.BatchNorm1d(128 * 7 * 7),

    Unflatten(N=-1, C=128, H=7, W=7),

    nn.ConvTranspose2d(128, 64, kernel_size=4, stride=2, padding=1),
    nn.ReLU(True),
    nn.BatchNorm2d(64),

    nn.ConvTranspose2d(64, 1, kernel_size=4, stride=2, padding=1),
    nn.Tanh()
)

```

```

# *****END OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****
#####
#                                     END OF YOUR CODE                                     #
#####

```

```
def run_a_gan(D, G, D_solver, G_solver, discriminator_loss, generator_loss, loader_train,
show_every=250,
```

```
    batch_size=128, noise_size=96, num_epochs=10):
```

```
    """
```

```
    Train a GAN!
```

```
    Inputs:
```

- D, G: PyTorch models for the discriminator and generator
- D\_solver, G\_solver: torch.optim Optimizers to use for training the discriminator and generator.
- discriminator\_loss, generator\_loss: Functions to use for computing the generator and discriminator loss, respectively.
- show\_every: Show samples after every show\_every iterations.
- batch\_size: Batch size to use for training.
- noise\_size: Dimension of the noise to use as input to the generator.

```

- num_epochs: Number of epochs over the training dataset to use for training.
"""
images = []
iter_count = 0
for epoch in range(num_epochs):
    for x, _ in loader_train:
        if len(x) != batch_size:
            continue
        D_solver.zero_grad()
        real_data = x.type(dtype)
        logits_real = D(2* (real_data - 0.5)).type(dtype)

        g_fake_seed = sample_noise(batch_size, noise_size).type(dtype)
        fake_images = G(g_fake_seed).detach()
        logits_fake = D(fake_images.view(batch_size, 1, 28, 28))

        d_total_error = discriminator_loss(logits_real, logits_fake)
        d_total_error.backward()
        D_solver.step()

        G_solver.zero_grad()
        g_fake_seed = sample_noise(batch_size, noise_size).type(dtype)
        fake_images = G(g_fake_seed)

        gen_logits_fake = D(fake_images.view(batch_size, 1, 28, 28))
        g_error = generator_loss(gen_logits_fake)
        g_error.backward()
        G_solver.step()

        if (iter_count % show_every == 0):
            print('Iter: {}, D: {:.4}, G:
{:.4}'.format(iter_count, d_total_error.item(), g_error.item()))
            imgs_numpy = fake_images.data.cpu().numpy()
            images.append(imgs_numpy[0:16])

    iter_count += 1

return images

```

```

class ChunkSampler(sampler.Sampler):
    """Samples elements sequentially from some offset.
    Arguments:
        num_samples: # of desired datapoints
        start: offset where we should start selecting from
    """
    def __init__(self, num_samples, start=0):
        self.num_samples = num_samples
        self.start = start

    def __iter__(self):
        return iter(range(self.start, self.start + self.num_samples))

    def __len__(self):
        return self.num_samples

```

```

class Flatten(nn.Module):
    def forward(self, x):
        N, C, H, W = x.size() # read in N, C, H, W
        return x.view(N, -1) # "flatten" the C * H * W values into a single vector per image

```

```

class Unflatten(nn.Module):
    """
    An Unflatten module receives an input of shape (N, C*H*W) and reshapes it
    to produce an output of shape (N, C, H, W).
    """

```

```

"""
def __init__(self, N=-1, C=128, H=7, W=7):
    super(Unflatten, self).__init__()
    self.N = N
    self.C = C
    self.H = H
    self.W = W
def forward(self, x):
    return x.view(self.N, self.C, self.H, self.W)

def initialize_weights(m):
    if isinstance(m, nn.Linear) or isinstance(m, nn.ConvTranspose2d):
        nn.init.xavier_uniform_(m.weight.data)

def preprocess_img(x):
    return 2 * x - 1.0

def deprocess_img(x):
    return (x + 1.0) / 2.0

def rel_error(x, y):
    return np.max(np.abs(x - y) / (np.maximum(1e-8, np.abs(x) + np.abs(y))))

def count_params(model):
    """Count the number of parameters in the model. """
    param_count = np.sum([np.prod(p.size()) for p in model.parameters()])
    return param_count

```