

```

import torch
import torch.nn as nn
import torch.nn.functional as F
from ResUNet import ConditionalUnet
from utils import *

device = torch.device("cuda" if torch.cuda.is_available() else "cpu")

class ConditionalDDPM(nn.Module):
    def __init__(self, dmconfig):
        super().__init__()
        self.dmconfig = dmconfig
        self.loss_fn = nn.MSELoss()
        self.network = ConditionalUnet(1, self.dmconfig.num_feat, self.dmconfig.num_classes)

    def scheduler(self, t_s):
        beta_1, beta_T, T = self.dmconfig.beta_1, self.dmconfig.beta_T, self.dmconfig.T
        # ===== #
        # YOUR CODE HERE:
        # Inputs:
        #     t_s: the input time steps, with shape (B,1).
        # Outputs:
        #     one dictionary containing the variance schedule
        #     $\beta_t$ along with other potentially useful constants.

        # Make sure t_s is shape (B,1)
        if t_s.dim() == 0:
            t_s = t_s.view(1,1)
        elif t_s.dim() == 1:
            t_s = t_s.unsqueeze(-1)

        B = t_s.size(0)    # number of examples in the batch

        # Build the scheduler for the DDPM process
        device = t_s.device
        betas = torch.linspace(beta_1, beta_T, T, device=device)
        alphas = 1.0 - betas
        alphaBars = torch.cumprod(alphas, dim=0)

        # Zero-base indices
        t = t_s.squeeze(-1).long() - 1
        t = t.clamp(0, T-1)

        # Based on spec
        beta_t = betas[t].view(B,1,1,1)
        sqrt_beta_t = torch.sqrt(beta_t)
        alpha_t = alphas[t].view(B,1,1,1)
        oneover_sqrt_alpha = 1.0 / torch.sqrt(alpha_t)
        alpha_t_bar = alphaBars[t].view(B,1,1,1)
        sqrt_alpha_bar = torch.sqrt(alpha_t_bar)
        sqrt_oneminus_alpha_bar = torch.sqrt(1.0 - alpha_t_bar)

        # ===== #
        return {
            'beta_t': beta_t,
            'sqrt_beta_t': sqrt_beta_t,
            'alpha_t': alpha_t,
            'sqrt_alpha_bar': sqrt_alpha_bar,
            'oneover_sqrt_alpha': oneover_sqrt_alpha,
            'alpha_t_bar': alpha_t_bar,
            'sqrt_oneminus_alpha_bar': sqrt_oneminus_alpha_bar
        }

    def forward(self, images, conditions):
        T = self.dmconfig.T
        noise_loss = None
        # ===== #

```

```

# YOUR CODE HERE:
# Complete the training forward process based on the
# given training algorithm.
# Inputs:
#     images: real images from the dataset, with size (B,1,28,28).
#     conditions: condition labels, with size (B). You should
#                 convert it to one-hot encoded labels with size (B,10)
#                 before making it as the input of the denoising network.
# Outputs:
#     noise_loss: loss computed by the self.loss_fn function .

B, _, H, W = images.shape
device = images.device

# One-hot encode the labels
c = F.one_hot(conditions, num_classes=self.dmconfig.num_classes).float() # (B,
num_classes)

# Classifier-free masking
p_uncond = getattr(self.dmconfig, 'p_uncond', 0.0)
if p_uncond > 0:
    mask = (torch.randn(B, device=device) < p_uncond)
    c[mask] = 0.0

# Sample for each batch element
t = torch.randint(1, T+1, (B,1), device=device)
eps = torch.randn_like(images)

# Compute the noisy input
sched = self.scheduler(t)

x_t = sched['sqrt_alpha_bar'] * images \
        + sched['sqrt_oneminus_alpha_bar'] * eps # broadcastâ (B,1,H,W)

# Predict the noise
noise_pred = self.network(x_t, t.squeeze(-1), c)

# MSE loss
noise_loss = self.loss_fn(noise_pred, eps)

# ===== #

return noise_loss

def sample(self, conditions, omega):
    T = self.dmconfig.T
    X_t = None
    # ===== #
    # YOUR CODE HERE:
    # Complete the training forward process based on the
    # given sampling algorithm.
    # Inputs:
    #     conditions: condition labels, with size (B). You should
    #                 convert it to one-hot encoded labels with size (B,10)
    #                 before making it as the input of the denoising network.
    #     omega: conditional guidance weight.
    # Outputs:
    #     generated_images

    B = conditions.size(0)
    device = conditions.device

    # One-hot encode labels
    c = F.one_hot(conditions, num_classes=self.dmconfig.num_classes).float()

    # Gaussian noise x_T
    X_t = torch.randn(B, 1, 28, 28, device=device)

```

```

# eval mode, no grad
self.network.eval()
with torch.no_grad():
    for t in range(T, 0, -1):
        # Build timestep tensors
        t_vec = torch.full((B,), t, device=device, dtype=torch.long)
        t_batch = t_vec.unsqueeze(-1)

        # Query scheduler
        sched = self.scheduler(t_batch)

        # Predict noise
        eps_cond = self.network(X_t, t_vec, c)
        eps_uncond = self.network(X_t, t_vec, torch.zeros_like(c))

        # Classifier-free guidance
        eps_hat = (1.0 + omega) * eps_cond - omega * eps_uncond

        # Compute the mean without noise
        coef = (1.0 - sched['alpha_t']) / sched['sqrt_oneminus_alpha_bar']
        X_prev = sched['oneover_sqrt_alpha'] * (X_t - coef * eps_hat)

        # Add noise if t > 1
        if t > 1:
            z = torch.randn_like(X_t)
            X_prev = X_prev + sched['sqrt_beta_t'] * z

        # Step down
        X_t = X_prev

# back to train mode
self.network.train()

# ===== #
generated_images = (X_t * 0.3081 + 0.1307).clamp(0,1) # denormalize the output images
return generated_images

# def sample(self, conditions, omega):
#     T = self.dmconfig.T
#     X_t = None
#     # ===== #
#     # YOUR CODE HERE:
#     # Complete the training forward process based on the
#     # given sampling algorithm.
#     # Inputs:
#     #     conditions: condition labels, with size (B). You should
#     #                 convert it to one-hot encoded labels with size (B,10)
#     #                 before making it as the input of the denoising network.
#     #     omega: conditional guidance weight.
#     # Outputs:
#     #     generated_images

#     B = conditions.size(0)
#     device = conditions.device

#     # One-hot encode
#     c = F.one_hot(conditions, num_classes=self.dmconfig.num_classes).float()

#     # Gaussian noise
#     X_t = torch.randn(B, 1, 28, 28, device=device)

#     # put the network into eval & disable grads
#     self.network.eval()
#     with torch.no_grad():
#         # Denoise

```

```

#         for t in range(T, 0, -1):
#             t_batch = torch.full((B,1), t, device=device, dtype=torch.long)
#             sched = self.scheduler(t_batch)

#             z = torch.randn_like(X_t) if t > 1 else torch.zeros_like(X_t)

#             eps_cond  = self.network(X_t, t_batch.squeeze(-1),          c)
#             eps_uncond = self.network(X_t, t_batch.squeeze(-1), torch.zeros_like(c))

#             # Classifier-free guidance
#             eps_hat = (1.0 + omega) * eps_cond - omega * eps_uncond

#             coef = (1.0 - sched['alpha_t']) / sched['sqrt_oneminus_alpha_bar']
#             X_t = (
#                 sched['oneover_sqrt_alpha']
#                 * (X_t - coef * eps_hat)
#                 + sched['sqrt_beta_t'] * z
#             )

#         # back to train mode
#         self.network.train()

#         # ===== #
#         generated_images = (X_t * 0.3081 + 0.1307).clamp(0,1) # denormalize the output
images
#         return generated_images

```