

```

import torch
import torch.optim as optim
import torch.nn.functional as F
import torch.nn
import gymnasium as gym
from replay_buffer import ReplayBufferDQN
import wandb
import random
import numpy as np
import os
import time
from utils import exponential_decay
import typing

```

*# TODO: change the logging here if you don't like wandb*

```
class DQN:
```

```

    def __init__(self, env:typing.Union[gym.Env,gym.Wrapper],
        #model params
        model:torch.nn.Module,
        model_kwargs:dict = {},
        #overall hyperparams
        lr:float = 0.001, gamma:float = 0.99,
        buffer_size:int = 10000, batch_size:int = 32,
        loss_fn:str = 'mse_loss',
        use_wandb:bool = False,
        device:str = 'cpu',
        seed:int = 42,
        epsilon_scheduler = exponential_decay(1,700,0.1),
        save_path:str = None):

```

*"""Initializes the DQN algorithm*

*Args:*

*env (gym.Env/gym.Wrapper): the environment to train on*  
*model (torch.nn.Module): the model to train*  
*model\_kwargs (dict, optional): the keyword arguments to pass to the model. Defaults*

*to {}.*

*lr (float, optional): the learning rate to use in the optimizer. Defaults to 0.001.*  
*gamma (float, optional): discount factor. Defaults to 0.99.*  
*buffer\_size (int, optional): the size of the replay buffer. Defaults to 10000.*  
*batch\_size (int, optional): the batch size. Defaults to 32.*  
*loss\_fn (str, optional): the name of the loss function to use. Defaults to*

*'mse\_loss'.*

*use\_wandb (bool, optional): \_description\_. Defaults to False.*  
*device (str, optional): \_description\_. Defaults to 'cpu'.*  
*seed (int, optional): the seed to use for reproducibility. Defaults to 42.*  
*epsilon\_scheduler ([type], optional): the epsilon scheduler to use, must have a*

*\_\_call\_\_ method that returns a float between 0 and 1*

*save\_path (str, optional): \_description\_. Defaults to None.*

*Raises:*

*ValueError: \_description\_*

*"""*

```
self.env = env
```

```
self._set_seed(seed)
```

```
self.observation_space = self.env.observation_space.shape
```

```
self.model = model(
```

```
    self.observation_space,
```

```
    self.env.action_space.n, **model_kwargs
```

```
) .to(device)
```

```
self.model.train()
```

```
self.optimizer = optim.Adam(self.model.parameters(), lr = lr)
```

```
self.gamma = gamma
```

```
self.replay_buffer = ReplayBufferDQN(buffer_size)
```

```

self.batch_size = batch_size
self.i_update = 0
self.device = device
self.epsilon_decay = epsilon_scheduler
self.save_path = save_path if save_path is not None else "./"

#set the loss function
if loss_fn == 'smooth_l1_loss':
    self.loss_fn = F.smooth_l1_loss
elif loss_fn == 'mse_loss':
    self.loss_fn = F.mse_loss
else:
    raise ValueError('loss_fn must be either smooth_l1_loss or mse_loss')

self.wandb = use_wandb
if self.wandb:
    wandb.init(project = 'racing-car-dqn')
    #log the hyperparameters
    wandb.config.update({
        'lr': lr,
        'gamma': gamma,
        'buffer_size': buffer_size,
        'batch_size': batch_size,
        'loss_fn': loss_fn,
        'device': device,
        'seed': seed,
        'save_path': save_path
    })

def train(self, n_episodes:int = 1000,validate_every:int = 100, n_validation_episodes:int
= 10, n_test_episodes:int = 10,save_every:int = 100):
    os.makedirs(self.save_path, exist_ok = True)
    best_val_reward = -np.inf

    for episode in range(n_episodes):
        state,_ = self.env.reset()
        done = False
        truncated = False
        total_reward = 0
        i = 0
        loss = 0
        start_time = time.time()
        epsilon = self.epsilon_decay()
        while (not done) and (not truncated):
            action = self._sample_action(state, epsilon)
            next_state, reward, done, truncated, _ = self.env.step(action)
            self.replay_buffer.add(state, action, reward, next_state, done)
            total_reward += reward
            state = next_state

            not_warm_starting,l = self._optimize_model()
            if not_warm_starting:
                loss += l
                epsilon = self.epsilon_decay()
                i += 1

        if i == 0:
            avg_loss = loss / (i+1)
        else:
            avg_loss = loss/i
        if self.wandb:
            wandb.log({'total_reward': total_reward, 'loss': avg_loss})
        print(f"Episode: {episode}: Time: {time.time() - start_time} Total Reward:
{total_reward} Avg_Loss: {avg_loss}")
        if episode % validate_every == validate_every - 1:
            mean_reward, std_reward = self.validate(n_validation_episodes)
            if self.wandb:
                wandb.log({'mean_reward': mean_reward, 'std_reward': std_reward})

```

```

        print("Validation Mean Reward: {}".format(mean_reward))
        print("Validation Std Reward: {}".format(std_reward))
        if mean_reward > best_val_reward:
            best_val_reward = mean_reward
            self._save('best')

    if episode % save_every == save_every - 1:
        self._save(str(episode))

    self._save('final')
    self.load_model('best')
    mean_reward, std_reward = self.validate(n_test_episodes)
    if self.wandb:
        wandb.log({'mean_test_reward': mean_reward, 'std_test_reward': std_reward})
    print("Test Mean Reward: {} Test Std Reward: {}".format(mean_reward, std_reward))

def _optimize_model(self):
    """
    Performs one optimization step on the DQN.

    Returns:
        bool: whether we have enough samples to optimize the model, which we define as
        having at least 10*batch_size samples
        float: the loss, if we do not have enough samples, we return 0
    """
    # ===== YOUR CODE HERE =====
    # TODO: some hints to get you started:
    # 1. check if the replay buffer has enough samples
    # 2. sample a minibatch
    # 3. No-op
    # 4. compute current Q: q_values = ...
    # 5. compute target Q
    # 6. compute loss between current Q and target Q
    # 7. backprop
    # =====

    # step 1
    if len(self.replay_buffer) < (10 * self.batch_size):
        return False, 0.0

    # step 2
    states, actions, rewards, next_states, dones =
self.replay_buffer.sample(self.batch_size, device=self.device)

    # step 3
    # No-op

    # step 4
    # Select Q-values for the taken actions
    q_values = self.model(states) # shape: (batch_size, num_actions)

    # step 5: target_q
    with torch.no_grad():
        next_q_values = self.model(next_states)
        max_next_q_values = next_q_values.max(dim=1)[0]
        if dones.all():
            targets = rewards
        else:
            targets = rewards + self.gamma * max_next_q_values

    # step 6
    q_values_given_action = q_values.gather(1, actions.unsqueeze(1)).squeeze(1) # shape:
(batch_size,)
    loss = self.loss_fn(q_values_given_action, targets)

    # step 7
    self.optimizer.zero_grad()

```

```

loss.backward()
self.optimizer.step()

return True, loss.item()

# ===== YOUR CODE ENDS =====

def _sample_action(self, state:np.ndarray
                  , epsilon:float = 0.1)->int:
    """
    Samples an action from the model

    Args:
        state (np.ndarray): the state, of shape [n_c,h,w]
        epsilon (float, optional): the epsilon for epsilon greedy. Defaults to 0.1.

    Returns:
        int: the index of the action to take
    """
    # ===== YOUR CODE HERE =====
    # TODO:
    # Epsilon-greedy action selection:
    # - if probability epsilon: random action
    # - else: greedy action
    # =====

    if random.random() < epsilon:
        index = self.env.action_space.sample()
    else:
        state_tensor = torch.tensor(state, dtype=torch.float32,
device=self.device).unsqueeze(0)
        with torch.no_grad():
            q_values = self.model(state_tensor)
            index = q_values.argmax(dim=1).item()

    # ===== YOUR CODE ENDS =====
    return index

def _set_seed(self, seed:int):
    random.seed(seed)
    np.random.seed(seed)
    self.seed = seed
    torch.manual_seed(seed)
    torch.cuda.manual_seed(seed)
    torch.backends.cudnn.deterministic = True
    gym.utils.seeding.np_random(seed)

def _validate_once(self):
    state,_ = self.env.reset()
    done = False
    truncated = False
    total_reward = 0
    i = 0
    # epsilon = self.epsilon_decay()
    while (not done) and (not truncated):
        action = self._sample_action(state, 0)
        # out = self.env.step(action)
        next_state, reward, done, truncated, _ = self.env.step(action)
        # next_state = np.array(state_buffer[-self.n_frames:])
        total_reward += reward
        state = next_state
    return total_reward

def validate(self, n_episodes:int = 10):
    # self.model.eval()
    rewards_per_episode = []

```

```

for _ in range(n_episodes):
    rewards_per_episode.append(self._validate_once())
    # self.model.train()
    return np.mean(rewards_per_episode), np.std(rewards_per_episode)

def load_model(self, suffix:str = ''):
    self.model.load_state_dict(torch.load(os.path.join(self.save_path,
f'model_{suffix}.pt'))))

def _save(self, suffix:str = ''):
    torch.save(self.model.state_dict(), os.path.join(self.save_path,
f'model_{suffix}.pt'))

def play_episode(self, epsilon:float = 0, return_frames:bool = True, seed:int = None):
    """Plays an episode of the environment

    Args:
        epsilon (float, optional): the epsilon for epsilon greedy. Defaults to 0.
        return_frames (bool, optional): whether we should return frames. Defaults to True.
        seed (int, optional): the seed for the enviroment. Defaults to None.

    Returns:
        if return_frames is True, returns the total reward and the frames
        if return_frames is False, returns the total reward
    """
    if seed is not None:
        state,_ = self.env.reset(seed = seed)
    else:
        state,_ = self.env.reset()

    done = False
    total_reward = 0
    if return_frames:
        frames = []

    with torch.no_grad():
        while not done:
            action = self._sample_action(state, epsilon)
            next_state, reward, terminated, truncated, _ = self.env.step(action)
            total_reward += reward
            done = terminated or truncated
            if return_frames:
                frames.append(self.env.render())
            state = next_state

    if return_frames:
        return total_reward, frames

    return total_reward

```

```

class HardUpdateDQN(DQN):

```

```

def __init__(self, env, model, model_kwargs:dict = {},
            update_freq:int = 5, *args, **kwargs):
    super().__init__(env, model, model_kwargs, *args, **kwargs)
    # ===== YOUR CODE HERE =====
    # TODO:
    # fill in the initialization and synchronization of the target model weights
    # =====

    # Initialize target network with same architecture
    self.target_model = model(
        self.observation_space,
        self.env.action_space.n,
        **model_kwargs

```

```

).to(self.device)

# Copy initial weights from model to target_model
self.target_model.load_state_dict(self.model.state_dict())

# Set target network to eval mode (not training)
self.target_model.eval()

# Store update frequency for hard target updates
self.update_freq = update_freq

# ===== YOUR CODE ENDS =====

def _optimize_model(self):
    """Optimizes the model

    Returns:
        bool: whether we have enough samples to optimize the model, which we define as
        having at least 10*batch_size samples
        float: the loss, if we do not have enough samples, we return 0
    """
    # ===== YOUR CODE HERE =====
    # TODO:
    # hint: you can copy over most of the code from the parent class
    # and only change one line
    # =====

    if len(self.replay_buffer) < (10 * self.batch_size):
        return False, 0.0

    # Sample batch
    states, actions, rewards, next_states, dones =
self.replay_buffer.sample(self.batch_size, device=self.device)

    # Compute Q(s, a) from current model
    q_values = self.model(states)
    q_values_given_action = q_values.gather(1, actions.unsqueeze(1)).squeeze(1)

    # Compute target Q using target network
    with torch.no_grad():
        next_q_values = self.target_model(next_states) # changed from self.model
        max_next_q_values = next_q_values.max(dim=1)[0]
        targets = rewards + self.gamma * max_next_q_values * (~dones)

    # Compute loss and update
    loss = self.loss_fn(q_values_given_action, targets)
    self.optimizer.zero_grad()
    loss.backward()
    self.optimizer.step()

    # Update target model if needed
    self._update_model()

    return True, loss.item()

# ===== YOUR CODE ENDS =====

def _update_model(self):
    self.i_update += 1
    if self.i_update % self.update_freq == 0:
        self.target_model.load_state_dict(self.model.state_dict())

def _save(self, suffix:str = ''):
    torch.save(self.model.state_dict(), os.path.join(self.save_path,
f'model_{suffix}.pt'))
    torch.save(self.target_model.state_dict(), os.path.join(self.save_path,
f'target_model_{suffix}.pt'))

```

```

def load_model(self, suffix:str = ''):
    self.model.load_state_dict(torch.load(os.path.join(self.save_path,
f'model_{suffix}.pt'))))
    self.target_model.load_state_dict(torch.load(os.path.join(self.save_path,
f'target_model_{suffix}.pt'))))

class SoftUpdateDQN(HardUpdateDQN):
    def __init__(self,env,model,model_kwargs:dict = {},
        tau:float = 0.01,*args,**kwargs):
        super().__init__(env,model,model_kwargs,*args,**kwargs)
        self.tau = tau

    def _update_model(self):
        """
        Soft updates the target model
        """
        # ===== YOUR CODE HERE =====
        # TODO
        # =====

        for target_param, model_param in zip(self.target_model.parameters(),
self.model.parameters()):
            target_param.data.copy_(self.tau * model_param.data + (1.0 - self.tau) *
target_param.data)

        # ===== YOUR CODE ENDS =====

```